

An Introduction to Deep Learning With Python

[2.1] A first look at a Neural Network

Prof. Yuzo Iano

pgs: 27 - 30

```
In [1]: import keras
        from keras.datasets import mnist
```

Using TensorFlow backend.

```
In [2]: (train_images, train_labels), (test_images, test_labels) = mnist.load_data()
```

```
In [3]: # Let's Look at the Training Data
        print('Train Shape', train_images.shape)
        print('Train Label Length', len(train_labels))
        train_labels
```

Train Shape (60000, 28, 28)

Train Label Length 60000

```
Out[3]: array([5, 0, 4, ..., 5, 6, 8], dtype=uint8)
```

```
In [4]: # The Test Data
        print('Test Shape', test_images.shape)
        print('Test Label Length', len(test_labels))
        test_labels
```

Test Shape (10000, 28, 28)

Test Label Length 10000

```
Out[4]: array([7, 2, 1, ..., 4, 5, 6], dtype=uint8)
```

The Network Architecture

```
In [5]: from keras import models
        from keras import layers

        network = models.Sequential()
        network.add(layers.Dense(512, activation = 'relu', input_shape = (28*28, )))
        network.add(layers.Dense(10, activation = 'softmax'))

        network.summary()
```

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 512)	401920
dense_2 (Dense)	(None, 10)	5130
Total params: 407,050		
Trainable params: 407,050		
Non-trainable params: 0		

The Compilation Step

```
In [6]: network.compile(optimizer = 'rmsprop',  
                        loss = 'categorical_crossentropy',  
                        metrics = ['accuracy'])
```

Preparing the image data

```
In [7]: train_images = train_images.reshape((60000, 28 * 28))  
train_images = train_images.astype('float32')/ 255  
  
test_images = test_images.reshape((10000, 28 * 28))  
test_images = test_images.astype('float32')/ 255
```

Preparing the labels

```
In [8]: from keras.utils import to_categorical  
  
train_labels = to_categorical(train_labels)  
test_labels = to_categorical(test_labels)
```

Train the model

```
In [9]: network.fit(train_images, train_labels, epochs = 5, batch_size = 128)  
  
Epoch 1/5  
60000/60000 [=====] - 9s 150us/step - loss: 0.2560 - acc: 0.  
9260  
Epoch 2/5  
60000/60000 [=====] - 5s 90us/step - loss: 0.1037 - acc: 0.9  
692  
Epoch 3/5  
60000/60000 [=====] - 5s 86us/step - loss: 0.0680 - acc: 0.9  
797  
Epoch 4/5  
60000/60000 [=====] - 5s 88us/step - loss: 0.0498 - acc: 0.9  
848  
Epoch 5/5  
60000/60000 [=====] - 5s 86us/step - loss: 0.0372 - acc: 0.9  
891
```

```
Out[9]: <keras.callbacks.History at 0x289ca49fa20>
```

```
In [10]: test_loss, test_acc = network.evaluate(test_images, test_labels)  
print('TEST ACC: ', test_acc)  
  
10000/10000 [=====] - 1s 56us/step  
TEST ACC: 0.98
```

Pablo Minango

- pablodavid218@gmail.com