

# An Introduction to Deep Learning With Python

## [3.2] Classifying movie reviews a binary classification example

Prof. Yuzo Iano

pgs: 68 - 77

### Loading the IMDB dataset

```
In [1]: from keras.datasets import imdb
(train_data, train_labels), (test_data, test_labels) = imdb.load_data(num_words = 10000)
```

Using TensorFlow backend.

```
In [2]: train_data[0:10000:10000]
```

```
Out[2]: array([[list([1, 14, 22, 16, 43, 530, 973, 1622, 1385, 65, 458, 4468, 66, 3941, 4, 173, 36, 256, 5,
25, 100, 43, 838, 112, 50, 670, 2, 9, 35, 480, 284, 5, 150, 4, 172, 112, 167, 2, 336, 385, 39, 4,
172, 4536, 1111, 17, 546, 38, 13, 447, 4, 192, 50, 16, 6, 147, 2025, 19, 14, 22, 4, 1920, 4613, 46
9, 4, 22, 71, 87, 12, 16, 43, 530, 38, 76, 15, 13, 1247, 4, 22, 17, 515, 17, 12, 16, 626, 18, 2,
5, 62, 386, 12, 8, 316, 8, 106, 5, 4, 2223, 5244, 16, 480, 66, 3785, 33, 4, 130, 12, 16, 38, 619,
5, 25, 124, 51, 36, 135, 48, 25, 1415, 33, 6, 22, 12, 215, 28, 77, 52, 5, 14, 407, 16, 82, 2, 8,
4, 107, 117, 5952, 15, 256, 4, 2, 7, 3766, 5, 723, 36, 71, 43, 530, 476, 26, 400, 317, 46, 7, 4,
2, 1029, 13, 104, 88, 4, 381, 15, 297, 98, 32, 2071, 56, 26, 141, 6, 194, 7486, 18, 4, 226, 22, 2
1, 134, 476, 26, 480, 5, 144, 30, 5535, 18, 51, 36, 28, 224, 92, 25, 104, 4, 226, 65, 16, 38, 133
4, 88, 12, 16, 283, 5, 16, 4472, 113, 103, 32, 15, 16, 5345, 19, 178, 32])],
      dtype=object)
```

```
In [3]: train_labels[0]
```

```
Out[3]: 1
```

```
In [4]: max([max(sequence) for sequence in train_data])
```

```
Out[4]: 9999
```

```
In [5]: word_index = imdb.get_word_index()
reverse_word_index = dict(
    [(value, key) for (key, value) in word_index.items()])
decoded_review = ' '.join(
    [reverse_word_index.get(i - 3, '?') for i in train_data[0]])
```

### Preparing the data

Encoding the Integer sequences Into a binary matrix

```
In [6]: import numpy as np

def vectorize_sequences(sequences, dimension=10000):
    results = np.zeros((len(sequences), dimension))
    for i, sequence in enumerate(sequences):
        results[i, sequence] = 1.
    return results

x_train = vectorize_sequences(train_data)
x_test = vectorize_sequences(test_data)
```

```
In [7]: x_train[0]
```

```
Out[7]: array([0., 1., 1., ..., 0., 0., 0.])
```

```
In [8]: y_train = np.asarray(train_labels).astype('float32')
y_test = np.asarray(test_labels).astype('float32')
```

## Building your model

The model definition

```
In [9]: from keras.models import Sequential
from keras.layers import Dense

model = Sequential()
model.add(Dense(16, activation='relu', input_shape=(10000,)))
model.add(Dense(16, activation='relu'))
model.add(Dense(1, activation='sigmoid'))
model.summary()
```

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 16)	160016
dense_2 (Dense)	(None, 16)	272
dense_3 (Dense)	(None, 1)	17
Total params: 160,305		
Trainable params: 160,305		
Non-trainable params: 0		

## Compiling the model

```
In [10]: model.compile(optimizer = 'rmsprop',
                      loss = 'binary_crossentropy',
                      metrics = ['accuracy'])
```

## Configuring the optimizer

```
In [11]: from keras import optimizers

model.compile(optimizer = optimizers.RMSprop(lr=0.001),
              loss = 'binary_crossentropy',
              metrics = ['accuracy'])
```

## Using custom losses and metrics

```
In [12]: from keras import losses
from keras import metrics

model.compile(optimizer = optimizers.RMSprop(lr=0.001),
              loss = losses.binary_crossentropy,
              metrics = [metrics.binary_accuracy])
```

## Setting aside a validation set

```
In [13]: x_val = x_train[: 10000]
partial_x_train = x_train[10000:]

y_val = y_train[:10000]
partial_y_train = y_train[10000:]
```

## Training your model

```
In [14]: model.compile(optimizer = 'rmsprop',
                        loss = 'binary_crossentropy',
                        metrics = ['acc'])

history = model.fit(partial_x_train,
                    partial_y_train,
                    epochs = 20,
                    batch_size = 512,
                    validation_data = (x_val, y_val))
```

Train on 15000 samples, validate on 10000 samples

```
Epoch 1/20
15000/15000 [=====] - 7s 484us/step - loss: 0.4976 - acc: 0.7953 - val_loss: 0.3717 - val_acc: 0.8719
Epoch 2/20
15000/15000 [=====] - 5s 360us/step - loss: 0.2958 - acc: 0.9045 - val_loss: 0.2991 - val_acc: 0.8909
Epoch 3/20
15000/15000 [=====] - 8s 541us/step - loss: 0.2160 - acc: 0.9285 - val_loss: 0.3086 - val_acc: 0.8714
Epoch 4/20
15000/15000 [=====] - 6s 384us/step - loss: 0.1741 - acc: 0.9431 - val_loss: 0.2829 - val_acc: 0.8845
Epoch 5/20
15000/15000 [=====] - 4s 268us/step - loss: 0.1414 - acc: 0.9542 - val_loss: 0.2863 - val_acc: 0.8850
Epoch 6/20
15000/15000 [=====] - 5s 340us/step - loss: 0.1143 - acc: 0.9653 - val_loss: 0.3089 - val_acc: 0.8811
Epoch 7/20
15000/15000 [=====] - 6s 391us/step - loss: 0.0971 - acc: 0.9711 - val_loss: 0.3146 - val_acc: 0.8842
Epoch 8/20
15000/15000 [=====] - 5s 302us/step - loss: 0.0803 - acc: 0.9765 - val_loss: 0.3869 - val_acc: 0.8657
Epoch 9/20
15000/15000 [=====] - 3s 232us/step - loss: 0.0659 - acc: 0.9820 - val_loss: 0.3650 - val_acc: 0.8778
Epoch 10/20
15000/15000 [=====] - 4s 234us/step - loss: 0.0554 - acc: 0.9850 - val_loss: 0.3861 - val_acc: 0.8791
Epoch 11/20
15000/15000 [=====] - 5s 347us/step - loss: 0.0455 - acc: 0.9884 - val_loss: 0.4181 - val_acc: 0.8761
Epoch 12/20
15000/15000 [=====] - 8s 559us/step - loss: 0.0387 - acc: 0.9911 - val_loss: 0.4521 - val_acc: 0.8697
Epoch 13/20
15000/15000 [=====] - 5s 334us/step - loss: 0.0298 - acc: 0.9937 - val_loss: 0.4712 - val_acc: 0.8731
Epoch 14/20
15000/15000 [=====] - 6s 390us/step - loss: 0.0244 - acc: 0.9949 - val_loss: 0.5022 - val_acc: 0.8714
Epoch 15/20
15000/15000 [=====] - 5s 346us/step - loss: 0.0186 - acc: 0.9974 - val_loss: 0.5315 - val_acc: 0.8692
Epoch 16/20
15000/15000 [=====] - 6s 427us/step - loss: 0.0156 - acc: 0.9981 - val_loss: 0.5705 - val_acc: 0.8693
Epoch 17/20
15000/15000 [=====] - 6s 370us/step - loss: 0.0151 - acc: 0.9975 - val_loss: 0.6005 - val_acc: 0.8685
Epoch 18/20
15000/15000 [=====] - 7s 465us/step - loss: 0.0099 - acc: 0.9987 - val_loss: 0.6447 - val_acc: 0.8667
Epoch 19/20
15000/15000 [=====] - 6s 376us/step - loss: 0.0060 - acc: 0.9999 - val_loss: 0.7091 - val_acc: 0.8575
Epoch 20/20
15000/15000 [=====] - 4s 243us/step - loss: 0.0082 - acc: 0.9986 - val_loss: 0.6950 - val_acc: 0.8650
```

```
In [15]: history_dict = history.history
history_dict.keys()
```

```
Out[15]: dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
```

### Plotting the training and validation loss


```
In [17]: import matplotlib.pyplot as plt

loss_values = history_dict['loss']
val_loss_values = history_dict['val_loss']

epochs = range(1, len(loss_values) + 1)

plt.plot(epochs, loss_values, 'bo', label='Training loss')
plt.plot(epochs, val_loss_values, 'b', label='Validation loss')
plt.title('Training and Validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.grid()

plt.show()
```




### Plotting the training and validation accuracy

```
In [18]: acc = history_dict['acc']
val_acc = history_dict['val_acc']

plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation ')
plt.title('Training and Validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.grid()

plt.show()
```



## Retraining a model from scratch

```
In [20]: model = Sequential()
model.add(Dense(16, activation='relu', input_shape=(10000,)))
model.add(Dense(16, activation='relu'))
model.add(Dense(1, activation='sigmoid'))

model.compile(optimizer='rmsprop',
              loss='binary_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=4, batch_size=512)
```

Epoch 1/4  
25000/25000 [=====] - 7s 283us/step - loss: 0.4584 - acc: 0.8136  
Epoch 2/4  
25000/25000 [=====] - 6s 221us/step - loss: 0.2630 - acc: 0.9093  
Epoch 3/4  
25000/25000 [=====] - 4s 169us/step - loss: 0.2005 - acc: 0.9284  
Epoch 4/4  
25000/25000 [=====] - 4s 172us/step - loss: 0.1685 - acc: 0.9389

Out[20]: <keras.callbacks.History at 0x1db10695898>

```
In [21]: results = model.evaluate(x_test, y_test)
print('Results [loss, acc] = ', results)
```

25000/25000 [=====] - 8s 318us/step  
Results [loss, acc] = [0.2992388512086868, 0.88244]

## Using a trained network to generate predictions on new data

```
In [22]: model.predict(x_test)
```

```
Out[22]: array([[0.2116582 ],  
                [0.999501  ],  
                [0.9304853 ],  
                ...,  
                [0.14387219],  
                [0.1034008  ],  
                [0.70603096]], dtype=float32)
```

**Pablo Minango**

- pablodavid218@gmail.com