

# An Introduction to Deep Learning With Python

## [5.3] Using a pretrained convnet

Prof. Yuzo Iano

pgs: 143 - 158

### Instantiating the VGG16 convolutional base

```
In [1]: from keras.applications import VGG16

conv_base = VGG16(weights='imagenet', include_top=False, input_shape=(150, 150, 3))
conv_base.summary()
```

Using TensorFlow backend.

WARNING:tensorflow:From C:\Users\pablo\AppData\Roaming\Python\Python36\site-packages\tensorflow\python\framework\op\_def\_library.py:263: colocate\_with (from tensorflow.python.framework.ops) is deprecated and will be removed in a future version.

Instructions for updating:

Colocations handled automatically by placer.

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	(None, 150, 150, 3)	0
<hr/>		
block1_conv1 (Conv2D)	(None, 150, 150, 64)	1792
<hr/>		
block1_conv2 (Conv2D)	(None, 150, 150, 64)	36928
<hr/>		
block1_pool (MaxPooling2D)	(None, 75, 75, 64)	0
<hr/>		
block2_conv1 (Conv2D)	(None, 75, 75, 128)	73856
<hr/>		
block2_conv2 (Conv2D)	(None, 75, 75, 128)	147584
<hr/>		
block2_pool (MaxPooling2D)	(None, 37, 37, 128)	0
<hr/>		
block3_conv1 (Conv2D)	(None, 37, 37, 256)	295168
<hr/>		
block3_conv2 (Conv2D)	(None, 37, 37, 256)	590080
<hr/>		
block3_conv3 (Conv2D)	(None, 37, 37, 256)	590080
<hr/>		
block3_pool (MaxPooling2D)	(None, 18, 18, 256)	0
<hr/>		
block4_conv1 (Conv2D)	(None, 18, 18, 512)	1180160
<hr/>		
block4_conv2 (Conv2D)	(None, 18, 18, 512)	2359808
<hr/>		
block4_conv3 (Conv2D)	(None, 18, 18, 512)	2359808
<hr/>		
block4_pool (MaxPooling2D)	(None, 9, 9, 512)	0
<hr/>		
block5_conv1 (Conv2D)	(None, 9, 9, 512)	2359808
<hr/>		
block5_conv2 (Conv2D)	(None, 9, 9, 512)	2359808
<hr/>		
block5_conv3 (Conv2D)	(None, 9, 9, 512)	2359808
<hr/>		
block5_pool (MaxPooling2D)	(None, 4, 4, 512)	0
=====		
Total params: 14,714,688		
Trainable params: 14,714,688		
Non-trainable params: 0		

## Extracting features using the pretrained convolutional base

```
In [2]: import os
import numpy as np
from keras.preprocessing.image import ImageDataGenerator

base_dir = '../CAP_5/cats_and_dogs_small'
train_dir = os.path.join(base_dir, 'train')
validation_dir = os.path.join(base_dir, 'validation')
test_dir = os.path.join(base_dir, 'test')

datagen = ImageDataGenerator(rescale=1./255)
batch_size = 20

def extract_features(directory, sample_count):
    features = np.zeros(shape=(sample_count, 4, 4, 512))
    labels = np.zeros(shape=(sample_count))
    generator = datagen.flow_from_directory(directory,
                                            target_size=(150, 150),
                                            batch_size=batch_size,
                                            class_mode='binary')

    i = 0
    for inputs_batch, labels_batch in generator:
        features_batch = conv_base.predict(inputs_batch)
        features[i * batch_size : (i + 1) * batch_size] = features_batch
        labels[i * batch_size : (i + 1) * batch_size] = labels_batch
        i += 1
    if i * batch_size >= sample_count:
        break
    return features, labels
```

```
In [3]: train_features, train_labels = extract_features(train_dir, 2000)
validation_features, validation_labels = extract_features(validation_dir, 1000)
test_features, test_labels = extract_features(test_dir, 1000)
```

Found 2000 images belonging to 2 classes.  
Found 1000 images belonging to 2 classes.  
Found 1000 images belonging to 2 classes.

```
In [4]: train_features = np.reshape(train_features, (2000, 4 * 4 * 512))
validation_features = np.reshape(validation_features, (1000, 4 * 4 * 512))
test_features = np.reshape(test_features, (1000, 4 * 4 * 512))
```

## Defining and training the densely connected classifier

```
In [5]: from keras.models import Sequential
from keras.layers import Dense, Dropout
from keras.optimizers import RMSprop
```

```
model = Sequential()
model.add(Dense(256, activation='relu', input_dim=4 * 4 * 512))
model.add(Dropout(0.5))
model.add(Dense(1, activation='sigmoid'))
model.summary()
```

WARNING:tensorflow:From C:\Users\pablo\Python\envs\DAVID\lib\site-packages\keras\backend\tensorflow\_backend.py:3445: calling dropout (from tensorflow.python.ops.nn\_ops) with keep\_prob is deprecated and will be removed in a future version.

Instructions for updating:

Please use `rate` instead of `keep\_prob`. Rate should be set to `rate = 1 - keep\_prob`.

Layer (type)	Output Shape	Param #
=====		
dense_1 (Dense)	(None, 256)	2097408
-----		
dropout_1 (Dropout)	(None, 256)	0
-----		
dense_2 (Dense)	(None, 1)	257
=====		
Total params: 2,097,665		
Trainable params: 2,097,665		
Non-trainable params: 0		
-----		

```
In [6]: model.compile(optimizer=RMSprop(lr=2e-5),  
                      loss='binary_crossentropy',  
                      metrics=['acc'])  
history = model.fit(train_features, train_labels,  
                    epochs=30,  
                    batch_size=20,  
                    validation_data=(validation_features, validation_labels))
```

WARNING:tensorflow:From C:\Users\pablo\AppData\Roaming\Python\Python36\site-packages\tensorflow\python\ops\math\_ops.py:3066: to\_int32 (from tensorflow.python.ops.math\_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.cast instead.

Train on 2000 samples, validate on 1000 samples

Epoch 1/30

2000/2000 [=====] - 3s 2ms/step - loss: 0.6150 - acc: 0.6705 - val\_loss: 0.4784 - val\_acc: 0.7820

Epoch 2/30

2000/2000 [=====] - 3s 1ms/step - loss: 0.4401 - acc: 0.8015 - val\_loss: 0.3699 - val\_acc: 0.8600

Epoch 3/30

2000/2000 [=====] - 3s 1ms/step - loss: 0.3712 - acc: 0.8405 - val\_loss: 0.3285 - val\_acc: 0.8750

Epoch 4/30

2000/2000 [=====] - 3s 1ms/step - loss: 0.3263 - acc: 0.8660 - val\_loss: 0.3045 - val\_acc: 0.8850

Epoch 5/30

2000/2000 [=====] - 3s 1ms/step - loss: 0.2912 - acc: 0.8885 - val\_loss: 0.2866 - val\_acc: 0.8930

Epoch 6/30

2000/2000 [=====] - 3s 1ms/step - loss: 0.2659 - acc: 0.8970 - val\_loss: 0.2792 - val\_acc: 0.8910

Epoch 7/30

2000/2000 [=====] - 3s 1ms/step - loss: 0.2523 - acc: 0.8975 - val\_loss: 0.2681 - val\_acc: 0.8920

Epoch 8/30

2000/2000 [=====] - 3s 1ms/step - loss: 0.2379 - acc: 0.9060 - val\_loss: 0.2629 - val\_acc: 0.8950

Epoch 9/30

2000/2000 [=====] - 3s 1ms/step - loss: 0.2242 - acc: 0.9165 - val\_loss: 0.2674 - val\_acc: 0.8900

Epoch 10/30

2000/2000 [=====] - 3s 1ms/step - loss: 0.2145 - acc: 0.9125 - val\_loss: 0.2505 - val\_acc: 0.8980

Epoch 11/30

2000/2000 [=====] - 3s 1ms/step - loss: 0.1948 - acc: 0.9295 - val\_loss: 0.2485 - val\_acc: 0.8990

Epoch 12/30

2000/2000 [=====] - 3s 1ms/step - loss: 0.1893 - acc: 0.9240 - val\_loss: 0.2450 - val\_acc: 0.9000

Epoch 13/30

2000/2000 [=====] - 3s 1ms/step - loss: 0.1788 - acc: 0.9315 - val\_loss: 0.2477 - val\_acc: 0.8940

Epoch 14/30

2000/2000 [=====] - 3s 1ms/step - loss: 0.1711 - acc: 0.9415 - val\_loss: 0.2428 - val\_acc: 0.9010

Epoch 15/30

2000/2000 [=====] - 3s 1ms/step - loss: 0.1666 - acc: 0.9375 - val\_loss: 0.2422 - val\_acc: 0.8980

Epoch 16/30

2000/2000 [=====] - 3s 1ms/step - loss: 0.1619 - acc: 0.9420 - val\_loss: 0.2433 - val\_acc: 0.8990

Epoch 17/30

2000/2000 [=====] - 3s 1ms/step - loss: 0.1574 - acc: 0.9445 - val\_loss: 0.2373 - val\_acc: 0.9040

Epoch 18/30

2000/2000 [=====] - 3s 1ms/step - loss: 0.1492 - acc: 0.9470 - val\_loss: 0.2368 - val\_acc: 0.9030

Epoch 19/30

2000/2000 [=====] - 3s 1ms/step - loss: 0.1337 - acc: 0.9540 - val\_loss: 0.2350 - val\_acc: 0.9060

Epoch 20/30

2000/2000 [=====] - 3s 1ms/step - loss: 0.1353 - acc: 0.9525 - val\_loss: 0.2352 - val\_acc: 0.9010

Epoch 21/30

2000/2000 [=====] - 3s 1ms/step - loss: 0.1306 - acc: 0.9585 - val\_loss: 0.2356 - val\_acc: 0.9020

Epoch 22/30

2000/2000 [=====] - 3s 1ms/step - loss: 0.1208 - acc: 0.9605 - val\_loss: 0.2378 - val\_acc: 0.9020

Epoch 23/30

2000/2000 [=====] - 3s 1ms/step - loss: 0.1137 - acc: 0.9610 - val\_loss: 0.2355 - val\_acc: 0.9050

Epoch 24/30

```
2000/2000 [=====] - 3s 1ms/step - loss: 0.1124 - acc: 0.9650 - val_loss:
0.2411 - val_acc: 0.9030
Epoch 25/30
2000/2000 [=====] - 3s 1ms/step - loss: 0.1089 - acc: 0.9630 - val_loss:
0.2399 - val_acc: 0.9000
Epoch 26/30
2000/2000 [=====] - 3s 1ms/step - loss: 0.1042 - acc: 0.9680 - val_loss:
0.2358 - val_acc: 0.9050
Epoch 27/30
2000/2000 [=====] - 3s 1ms/step - loss: 0.0959 - acc: 0.9755 - val_loss:
0.2371 - val_acc: 0.9050
Epoch 28/30
2000/2000 [=====] - 3s 1ms/step - loss: 0.0928 - acc: 0.9720 - val_loss:
0.2408 - val_acc: 0.9050
Epoch 29/30
2000/2000 [=====] - 3s 1ms/step - loss: 0.0935 - acc: 0.9685 - val_loss:
0.2472 - val_acc: 0.9030
Epoch 30/30
2000/2000 [=====] - 3s 1ms/step - loss: 0.0878 - acc: 0.9740 - val_loss:
0.2393 - val_acc: 0.9030
```

## Plotting the results

```
In [21]: import matplotlib.pyplot as plt



acc = history.history['acc']
val_acc = history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']

epochs = range(1, len(acc) + 1)

plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.grid()
plt.legend()

plt.figure()

plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.grid()
plt.legend()
plt.show()
```



**Adding a densely connected classifier on top of the convolutional base**

```
In [8]: from keras.layers import Flatten

model = Sequential()
model.add(conv_base)
model.add(Flatten())
model.add(Dense(256, activation='relu'))
model.add(Dense(1, activation='sigmoid'))
model.summary()
```

Layer (type)	Output Shape	Param #
vgg16 (Model)	(None, 4, 4, 512)	14714688
flatten_1 (Flatten)	(None, 8192)	0
dense_3 (Dense)	(None, 256)	2097408
dense_4 (Dense)	(None, 1)	257
Total params: 16,812,353		
Trainable params: 16,812,353		
Non-trainable params: 0		

```
In [9]: print('This is the number of trainable weights before freezing the conv base: ', len(model.trainable_weights))
```

This is the number of trainable weights before freezing the conv base: 30

```
In [10]: conv_base.trainable = False
print('This is the number of trainable weights before freezing the conv base: ', len(model.trainable_weights))
```

This is the number of trainable weights before freezing the conv base: 4

## Training the model end to end with a frozen convolutional base

```
In [11]: from keras.preprocessing.image import ImageDataGenerator

train_datagen = ImageDataGenerator(rescale=1./255,
                                    rotation_range=40,
                                    width_shift_range=0.2,
                                    height_shift_range=0.2,
                                    shear_range=0.2,
                                    zoom_range=0.2,
                                    horizontal_flip=True,
                                    fill_mode='nearest')

test_datagen = ImageDataGenerator(rescale=1./255)

train_generator = train_datagen.flow_from_directory(train_dir,
                                                    target_size=(150, 150),
                                                    batch_size=20,
                                                    class_mode='binary')
validation_generator = test_datagen.flow_from_directory(validation_dir,
                                                        target_size=(150, 150),
                                                        batch_size=20,
                                                        class_mode='binary')
```

Found 2000 images belonging to 2 classes.  
Found 1000 images belonging to 2 classes.



```
In [12]: model.compile(loss='binary_crossentropy',  
                        optimizer=RMSprop(lr=2e-5),  
                        metrics=['acc'])  
history = model.fit_generator(train_generator,  
                              steps_per_epoch=100,  
                              epochs=30,  
                              validation_data=validation_generator,  
                              validation_steps=50)
```

Epoch 1/30  
100/100 [=====] - 308s 3s/step - loss: 0.5994 - acc: 0.6860 - val\_loss: 0.4581 - val\_acc: 0.8320  
Epoch 2/30  
100/100 [=====] - 307s 3s/step - loss: 0.5016 - acc: 0.7815 - val\_loss: 0.3908 - val\_acc: 0.8490  
Epoch 3/30  
100/100 [=====] - 306s 3s/step - loss: 0.4370 - acc: 0.8240 - val\_loss: 0.3407 - val\_acc: 0.8700  
Epoch 4/30  
100/100 [=====] - 308s 3s/step - loss: 0.4007 - acc: 0.8195 - val\_loss: 0.3150 - val\_acc: 0.8790  
Epoch 5/30  
100/100 [=====] - 305s 3s/step - loss: 0.3875 - acc: 0.8330 - val\_loss: 0.2976 - val\_acc: 0.8800  
Epoch 6/30  
100/100 [=====] - 305s 3s/step - loss: 0.3818 - acc: 0.8255 - val\_loss: 0.2861 - val\_acc: 0.8870  
Epoch 7/30  
100/100 [=====] - 309s 3s/step - loss: 0.3676 - acc: 0.8425 - val\_loss: 0.2909 - val\_acc: 0.8780  
Epoch 8/30  
100/100 [=====] - 305s 3s/step - loss: 0.3544 - acc: 0.8460 - val\_loss: 0.2716 - val\_acc: 0.8930  
Epoch 9/30  
100/100 [=====] - 309s 3s/step - loss: 0.3454 - acc: 0.8500 - val\_loss: 0.2664 - val\_acc: 0.9000  
Epoch 10/30  
100/100 [=====] - 307s 3s/step - loss: 0.3406 - acc: 0.8540 - val\_loss: 0.2643 - val\_acc: 0.8850  
Epoch 11/30  
100/100 [=====] - 307s 3s/step - loss: 0.3370 - acc: 0.8465 - val\_loss: 0.2666 - val\_acc: 0.8990  
Epoch 12/30  
100/100 [=====] - 308s 3s/step - loss: 0.3347 - acc: 0.8490 - val\_loss: 0.2554 - val\_acc: 0.8990  
Epoch 13/30  
100/100 [=====] - 306s 3s/step - loss: 0.3227 - acc: 0.8540 - val\_loss: 0.2506 - val\_acc: 0.9000  
Epoch 14/30  
100/100 [=====] - 308s 3s/step - loss: 0.3188 - acc: 0.8615 - val\_loss: 0.2836 - val\_acc: 0.8750  
Epoch 15/30  
100/100 [=====] - 305s 3s/step - loss: 0.3194 - acc: 0.8655 - val\_loss: 0.2525 - val\_acc: 0.8910  
Epoch 16/30  
100/100 [=====] - 309s 3s/step - loss: 0.3116 - acc: 0.8730 - val\_loss: 0.2601 - val\_acc: 0.8970  
Epoch 17/30  
100/100 [=====] - 306s 3s/step - loss: 0.3136 - acc: 0.8510 - val\_loss: 0.2521 - val\_acc: 0.8900  
Epoch 18/30  
100/100 [=====] - 309s 3s/step - loss: 0.3155 - acc: 0.8665 - val\_loss: 0.2490 - val\_acc: 0.8940  
Epoch 19/30  
100/100 [=====] - 307s 3s/step - loss: 0.3003 - acc: 0.8725 - val\_loss: 0.2461 - val\_acc: 0.8950  
Epoch 20/30  
100/100 [=====] - 307s 3s/step - loss: 0.3107 - acc: 0.8640 - val\_loss: 0.2598 - val\_acc: 0.8860  
Epoch 21/30  
100/100 [=====] - 308s 3s/step - loss: 0.3096 - acc: 0.8650 - val\_loss: 0.2455 - val\_acc: 0.8960  
Epoch 22/30  
100/100 [=====] - 306s 3s/step - loss: 0.3024 - acc: 0.8675 - val\_loss: 0.2419 - val\_acc: 0.9000  
Epoch 23/30  
100/100 [=====] - 309s 3s/step - loss: 0.3027 - acc: 0.8715 - val\_loss: 0.2464 - val\_acc: 0.9010  
Epoch 24/30  
100/100 [=====] - 307s 3s/step - loss: 0.2884 - acc: 0.8775 - val\_loss: 0.2634 - val\_acc: 0.8840  
Epoch 25/30  
100/100 [=====] - 309s 3s/step - loss: 0.2891 - acc: 0.8730 - val\_loss: 0.2399 - val\_acc: 0.9010  
Epoch 26/30

```
100/100 [=====] - 306s 3s/step - loss: 0.2955 - acc: 0.8680 - val_loss:
0.2382 - val_acc: 0.9000
Epoch 27/30
100/100 [=====] - 307s 3s/step - loss: 0.2865 - acc: 0.8760 - val_loss:
0.2375 - val_acc: 0.9040
Epoch 28/30
100/100 [=====] - 307s 3s/step - loss: 0.2886 - acc: 0.8745 - val_loss:
0.2547 - val_acc: 0.8910
Epoch 29/30
100/100 [=====] - 307s 3s/step - loss: 0.2914 - acc: 0.8745 - val_loss:
0.2395 - val_acc: 0.9020
Epoch 30/30
100/100 [=====] - 309s 3s/step - loss: 0.2888 - acc: 0.8820 - val_loss:
0.2396 - val_acc: 0.9050
```

## Plotting the results

```
In [13]: import matplotlib.pyplot as plt



acc = history.history['acc']
val_acc = history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']

epochs = range(1, len(acc) + 1)

plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.grid()
plt.legend()

plt.figure()

plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.grid()
plt.legend()
plt.show()
```



```
In [14]: conv_base.summary()
```

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	(None, 150, 150, 3)	0
block1_conv1 (Conv2D)	(None, 150, 150, 64)	1792
block1_conv2 (Conv2D)	(None, 150, 150, 64)	36928
block1_pool (MaxPooling2D)	(None, 75, 75, 64)	0
block2_conv1 (Conv2D)	(None, 75, 75, 128)	73856
block2_conv2 (Conv2D)	(None, 75, 75, 128)	147584
block2_pool (MaxPooling2D)	(None, 37, 37, 128)	0
block3_conv1 (Conv2D)	(None, 37, 37, 256)	295168
block3_conv2 (Conv2D)	(None, 37, 37, 256)	590080
block3_conv3 (Conv2D)	(None, 37, 37, 256)	590080
block3_pool (MaxPooling2D)	(None, 18, 18, 256)	0
block4_conv1 (Conv2D)	(None, 18, 18, 512)	1180160
block4_conv2 (Conv2D)	(None, 18, 18, 512)	2359808
block4_conv3 (Conv2D)	(None, 18, 18, 512)	2359808
block4_pool (MaxPooling2D)	(None, 9, 9, 512)	0
block5_conv1 (Conv2D)	(None, 9, 9, 512)	2359808
block5_conv2 (Conv2D)	(None, 9, 9, 512)	2359808
block5_conv3 (Conv2D)	(None, 9, 9, 512)	2359808
block5_pool (MaxPooling2D)	(None, 4, 4, 512)	0
=====		
Total params: 14,714,688		
Trainable params: 0		
Non-trainable params: 14,714,688		

### Freezing all layers up to a specific one

```
In [15]: conv_base.trainable = True

set_trainable = False
for layer in conv_base.layers:
    if layer.name == 'block5_conv1':
        set_trainable = True
    if set_trainable:
        layer.trainable = True
    else:
        layer.trainable = False
```

### Fine-tuning the model

```
In [16]: model.compile(loss='binary_crossentropy',
                        optimizer=RMSprop(lr=1e-5),
                        metrics=['acc'])

history = model.fit_generator(train_generator,
                              steps_per_epoch=100,
                              epochs=30,
                              validation_data=validation_generator,
                              validation_steps=50)
```

Epoch 1/30  
100/100 [=====] - 356s 4s/step - loss: 0.3044 - acc: 0.8570 - val\_loss:  
0.2197 - val\_acc: 0.9110  
Epoch 2/30  
100/100 [=====] - 356s 4s/step - loss: 0.2608 - acc: 0.8935 - val\_loss:  
0.2561 - val\_acc: 0.8960  
Epoch 3/30  
100/100 [=====] - 355s 4s/step - loss: 0.2253 - acc: 0.9130 - val\_loss:  
0.2110 - val\_acc: 0.9160  
Epoch 4/30  
100/100 [=====] - 355s 4s/step - loss: 0.2120 - acc: 0.9110 - val\_loss:  
0.2456 - val\_acc: 0.9100  
Epoch 5/30  
100/100 [=====] - 356s 4s/step - loss: 0.2135 - acc: 0.9105 - val\_loss:  
0.1803 - val\_acc: 0.9250  
Epoch 6/30  
100/100 [=====] - 355s 4s/step - loss: 0.1866 - acc: 0.9195 - val\_loss:  
0.2302 - val\_acc: 0.9080  
Epoch 7/30  
100/100 [=====] - 357s 4s/step - loss: 0.1800 - acc: 0.9250 - val\_loss:  
0.1749 - val\_acc: 0.9390  
Epoch 8/30  
100/100 [=====] - 358s 4s/step - loss: 0.1837 - acc: 0.9260 - val\_loss:  
0.2963 - val\_acc: 0.8860  
Epoch 9/30  
100/100 [=====] - 357s 4s/step - loss: 0.1669 - acc: 0.9365 - val\_loss:  
0.2177 - val\_acc: 0.9220  
Epoch 10/30  
100/100 [=====] - 356s 4s/step - loss: 0.1593 - acc: 0.9395 - val\_loss:  
0.1684 - val\_acc: 0.9310  
Epoch 11/30  
100/100 [=====] - 355s 4s/step - loss: 0.1452 - acc: 0.9410 - val\_loss:  
0.1851 - val\_acc: 0.9300  
Epoch 12/30  
100/100 [=====] - 355s 4s/step - loss: 0.1572 - acc: 0.9395 - val\_loss:  
0.1787 - val\_acc: 0.9220  
Epoch 13/30  
100/100 [=====] - 355s 4s/step - loss: 0.1310 - acc: 0.9510 - val\_loss:  
0.1785 - val\_acc: 0.9270  
Epoch 14/30  
100/100 [=====] - 355s 4s/step - loss: 0.1320 - acc: 0.9465 - val\_loss:  
0.1766 - val\_acc: 0.9360  
Epoch 15/30  
100/100 [=====] - 355s 4s/step - loss: 0.1155 - acc: 0.9540 - val\_loss:  
0.2552 - val\_acc: 0.9150  
Epoch 16/30  
100/100 [=====] - 355s 4s/step - loss: 0.1278 - acc: 0.9500 - val\_loss:  
0.2094 - val\_acc: 0.9270  
Epoch 17/30  
100/100 [=====] - 355s 4s/step - loss: 0.1175 - acc: 0.9535 - val\_loss:  
0.1324 - val\_acc: 0.9540  
Epoch 18/30  
100/100 [=====] - 355s 4s/step - loss: 0.1053 - acc: 0.9565 - val\_loss:  
0.2456 - val\_acc: 0.9230  
Epoch 19/30  
100/100 [=====] - 355s 4s/step - loss: 0.1138 - acc: 0.9585 - val\_loss:  
0.2398 - val\_acc: 0.9140  
Epoch 20/30  
100/100 [=====] - 356s 4s/step - loss: 0.1057 - acc: 0.9615 - val\_loss:  
0.1605 - val\_acc: 0.9450  
Epoch 21/30  
100/100 [=====] - 355s 4s/step - loss: 0.1159 - acc: 0.9580 - val\_loss:  
0.1927 - val\_acc: 0.9320  
Epoch 22/30  
100/100 [=====] - 356s 4s/step - loss: 0.0821 - acc: 0.9665 - val\_loss:  
0.2450 - val\_acc: 0.9230  
Epoch 23/30  
100/100 [=====] - 355s 4s/step - loss: 0.0976 - acc: 0.9670 - val\_loss:  
0.1997 - val\_acc: 0.9330  
Epoch 24/30  
100/100 [=====] - 356s 4s/step - loss: 0.0843 - acc: 0.9680 - val\_loss:  
0.2280 - val\_acc: 0.9260  
Epoch 25/30  
100/100 [=====] - 356s 4s/step - loss: 0.0796 - acc: 0.9720 - val\_loss:  
0.1918 - val\_acc: 0.9330  
Epoch 26/30

```
100/100 [=====] - 355s 4s/step - loss: 0.0792 - acc: 0.9720 - val_loss:
0.1751 - val_acc: 0.9410
Epoch 27/30
100/100 [=====] - 355s 4s/step - loss: 0.0761 - acc: 0.9680 - val_loss:
0.2232 - val_acc: 0.9350
Epoch 28/30
100/100 [=====] - 355s 4s/step - loss: 0.0779 - acc: 0.9680 - val_loss:
0.2875 - val_acc: 0.9250
Epoch 29/30
100/100 [=====] - 355s 4s/step - loss: 0.0807 - acc: 0.9700 - val_loss:
0.2408 - val_acc: 0.9240
Epoch 30/30
100/100 [=====] - 356s 4s/step - loss: 0.0664 - acc: 0.9745 - val_loss:
0.2850 - val_acc: 0.9190
```

## Plotting the results



```
In [17]: import matplotlib.pyplot as plt



acc = history.history['acc']
val_acc = history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']

epochs = range(1, len(acc) + 1)

plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.grid()
plt.legend()

plt.figure()

plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.grid()
plt.legend()
plt.show()
```





### Smoothing the plots

```
In [18]: def smooth_curve(points, factor=0.8):
    smoothed_points = []
    for point in points:
        if smoothed_points:
            previous = smoothed_points[-1]
            smoothed_points.append(previous * factor + point * (1 - factor))
        else:
            smoothed_points.append(point)
    return smoothed_points
```

```
In [19]: plt.plot(epochs, smooth_curve(acc), 'bo', label='Training acc')
plt.plot(epochs, smooth_curve(val_acc), 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.grid()
plt.legend()

plt.figure()

plt.plot(epochs, smooth_curve(loss), 'bo', label='Training loss')
plt.plot(epochs, smooth_curve(val_loss), 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.grid()
plt.legend()
plt.show()
```



```
In [20]: test_generator = test_datagen.flow_from_directory(test_dir,
                                                         target_size=(150, 150),
                                                         batch_size=20,
                                                         class_mode='binary')

test_loss, test_acc = model.evaluate_generator(test_generator, steps=50)
print('Test acc: ', test_acc)
```

Found 1000 images belonging to 2 classes.  
Test acc: 0.9249999916553497

**Pablo Minango**

- pablodavid218@gmail.com