

An Introduction to Deep Learning With Python

[6.3] Understanding recurrent neural networks

Prof. Yuzo Iano

pgs: 196 - 205

Pseudocode RNN

```
In [ ]: state_t = 0
        for input_t in input_sequence:
            output_t = f(input_t, state_t)
            state_t = output_t
```

More detailed pseudocode for the RNN

```
In [ ]: state_t = 0
        for input_t in input_sequence:
            output_t = activation(dot(W, input_t) + dot(U, state_t) + b)
            state_t = output_t
```

Numpy implementation of a simple RNN

```
In [1]: import numpy as np

timesteps = 100
input_features = 32
output_features = 64

inputs = np.random.random((timesteps, input_features))
state_t = np.zeros((output_features,))

W = np.random.random((output_features, input_features))
U = np.random.random((output_features, output_features))
b = np.random.random((output_features,))

successive_outputs = []

for input_t in inputs:
    output_t = np.tanh(np.dot(W, input_t) + np.dot(U, state_t) + b)
    successive_outputs.append(output_t)
    state_t = output_t

final_output_sequence = np.concatenate(successive_outputs, axis=0)
print(final_output_sequence)

[0.99999789 0.99977317 0.99999977 ... 1.          1.          1.          ]
```

A recurrent layer in Keras

```
In [2]: from keras.layers import SimpleRNN, Embedding
        from keras.models import Sequential
```

```
model = Sequential()
model.add(Embedding(10000, 2))
model.add(SimpleRNN(32))
model.summary()
```

Using TensorFlow backend.

WARNING:tensorflow:From C:\Users\pablo\AppData\Roaming\Python\Python36\site-packages\tensorflow\python\framework\op_def_library.py:263: colocate_with (from tensorflow.python.framework.ops) is deprecated and will be removed in a future version.

Instructions for updating:

Colocations handled automatically by placer.

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, None, 2)	20000
simple_rnn_1 (SimpleRNN)	(None, 32)	1120
Total params: 21,120		
Trainable params: 21,120		
Non-trainable params: 0		

```
In [3]: model = Sequential()
        model.add(Embedding(10000, 2))
        model.add(SimpleRNN(32, return_sequences=True))
        model.summary()
```

Layer (type)	Output Shape	Param #
embedding_2 (Embedding)	(None, None, 2)	20000
simple_rnn_2 (SimpleRNN)	(None, None, 32)	1120
Total params: 21,120		
Trainable params: 21,120		
Non-trainable params: 0		

```
In [4]: model = Sequential()
        model.add(Embedding(10000, 2))
        model.add(SimpleRNN(32, return_sequences=True))
        model.add(SimpleRNN(32, return_sequences=True))
        model.add(SimpleRNN(32, return_sequences=True))
        model.add(SimpleRNN(32))

        model.summary()
```

Layer (type)	Output Shape	Param #
embedding_3 (Embedding)	(None, None, 2)	20000
simple_rnn_3 (SimpleRNN)	(None, None, 32)	1120
simple_rnn_4 (SimpleRNN)	(None, None, 32)	2080
simple_rnn_5 (SimpleRNN)	(None, None, 32)	2080
simple_rnn_6 (SimpleRNN)	(None, 32)	2080
Total params: 27,360		
Trainable params: 27,360		
Non-trainable params: 0		

```
In [5]: from keras.datasets import imdb
        from keras.preprocessing import sequence

        max_features = 10000
        maxlen = 500
        batch_size = 32

        print('Loading data...')
        (input_train, y_train), (input_test, y_test) = imdb.load_data(num_words=max_features)
        print(len(input_train), 'train sequences')
        print(len(input_test), 'test sequences')

        print('Pad sequences (samples x time)')
        input_train = sequence.pad_sequences(input_train, maxlen=maxlen)
        input_test = sequence.pad_sequences(input_test, maxlen=maxlen)
        print('input_train shape: ', input_train.shape)
        print('input_test shape: ', input_test.shape)
```

```
Loading data...
25000 train sequences
25000 test sequences
Pad sequences (samples x time)
input_train shape: (25000, 500)
input_test shape: (25000, 500)
```

Training the model with Embedding and SimpleRNN layers

```
In [6]: from keras.layers import Dense

model = Sequential()
model.add(Embedding(max_features, 32))
model.add(SimpleRNN(32))
model.add(Dense(1, activation='sigmoid'))

model.compile(optimizer='rmsprop',
              loss='binary_crossentropy',
              metrics=['acc'])

history = model.fit(input_train, y_train,
                   epochs=10,
                   batch_size=128,
                   validation_split=0.2)
```

WARNING:tensorflow:From C:\Users\pablo\AppData\Roaming\Python\Python36\site-packages\tensorflow\python\ops\math_ops.py:3066: to_int32 (from tensorflow.python.ops.math_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.cast instead.

Train on 20000 samples, validate on 5000 samples

```
Epoch 1/10
20000/20000 [=====] - 15s 758us/step - loss: 0.6521 - acc: 0.6042 - val_loss: 0.6121 - val_acc: 0.6576
Epoch 2/10
20000/20000 [=====] - 15s 737us/step - loss: 0.4215 - acc: 0.8179 - val_loss: 0.4364 - val_acc: 0.8000
Epoch 3/10
20000/20000 [=====] - 15s 726us/step - loss: 0.3063 - acc: 0.8754 - val_loss: 0.3803 - val_acc: 0.8346
Epoch 4/10
20000/20000 [=====] - 14s 703us/step - loss: 0.2256 - acc: 0.9135 - val_loss: 0.4258 - val_acc: 0.8196
Epoch 5/10
20000/20000 [=====] - 14s 711us/step - loss: 0.1610 - acc: 0.9411 - val_loss: 0.4326 - val_acc: 0.8094
Epoch 6/10
20000/20000 [=====] - 14s 710us/step - loss: 0.0980 - acc: 0.9669 - val_loss: 0.4231 - val_acc: 0.8524
Epoch 7/10
20000/20000 [=====] - 14s 707us/step - loss: 0.0639 - acc: 0.9801 - val_loss: 0.4969 - val_acc: 0.8342
Epoch 8/10
20000/20000 [=====] - 15s 733us/step - loss: 0.0382 - acc: 0.9889 - val_loss: 0.5911 - val_acc: 0.8086
Epoch 9/10
20000/20000 [=====] - 15s 747us/step - loss: 0.0213 - acc: 0.9940 - val_loss: 0.6207 - val_acc: 0.8138
Epoch 10/10
20000/20000 [=====] - 15s 752us/step - loss: 0.0168 - acc: 0.9946 - val_loss: 0.7151 - val_acc: 0.7918
```

Plotting results

```
In [8]: import matplotlib.pyplot as plt


acc = history.history['acc']
val_acc = history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']

epochs = range(1, len(acc) + 1)

plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.grid()

plt.figure()

plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.grid()
plt.show()
```



Understanding the LSTM and GRU layers

Pseudocode details of the LSTM architecture (1/2)

```
In [ ]: output_t = activation(dot(state_t, Uo) + dot(input_t, Wo) + dot(C_t, Vo) + bo)
i_t = activation(dot(state_t, Ui) + dot(input_t, Wi) + bi)
f_t = activation(dot(state_t, Uf) + dot(input_t, Wf) + bf)
k_t = activation(dot(state_t, Uk) + dot(input_t, Wk) + bk)
```

Pseudocode details of the LSTM architecture (2/2)

```
In [ ]: c_t+1 = i_t * k_t + c_t * f_t
```

Using the LSTM layer in Keras

```
In [9]: from keras.layers import LSTM

model = Sequential()
model.add(Embedding(max_features, 32))
model.add(LSTM(32))
model.add(Dense(1, activation='sigmoid'))

model.compile(optimizer='rmsprop',
              loss='binary_crossentropy',
              metrics=['acc'])

history = model.fit(input_train, y_train,
                   epochs=10,
                   batch_size=128,
                   validation_split=0.2)
```

Train on 20000 samples, validate on 5000 samples

Epoch 1/10

20000/20000 [=====] - 54s 3ms/step - loss: 0.5075 - acc: 0.7637 - val_loss: 0.3946 - val_acc: 0.8276

Epoch 2/10

20000/20000 [=====] - 55s 3ms/step - loss: 0.2888 - acc: 0.8849 - val_loss: 0.3999 - val_acc: 0.8218

Epoch 3/10

20000/20000 [=====] - 51s 3ms/step - loss: 0.2317 - acc: 0.9120 - val_loss: 0.2969 - val_acc: 0.8884

Epoch 4/10

20000/20000 [=====] - 50s 3ms/step - loss: 0.1938 - acc: 0.9271 - val_loss: 0.3412 - val_acc: 0.8884

Epoch 5/10

20000/20000 [=====] - 54s 3ms/step - loss: 0.1746 - acc: 0.9360 - val_loss: 0.3060 - val_acc: 0.8804

Epoch 6/10

20000/20000 [=====] - 54s 3ms/step - loss: 0.1534 - acc: 0.9452 - val_loss: 0.4010 - val_acc: 0.8752

Epoch 7/10

20000/20000 [=====] - 53s 3ms/step - loss: 0.1404 - acc: 0.9506 - val_loss: 0.3010 - val_acc: 0.8872

Epoch 8/10

20000/20000 [=====] - 54s 3ms/step - loss: 0.1277 - acc: 0.9561 - val_loss: 0.3145 - val_acc: 0.8764

Epoch 9/10

20000/20000 [=====] - 57s 3ms/step - loss: 0.1185 - acc: 0.9572 - val_loss: 0.3592 - val_acc: 0.8720

Epoch 10/10

20000/20000 [=====] - 57s 3ms/step - loss: 0.1091 - acc: 0.9615 - val_loss: 0.4501 - val_acc: 0.8728

Plotting the results

```
In [10]: import matplotlib.pyplot as plt



acc = history.history['acc']
val_acc = history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']

epochs = range(1, len(acc) + 1)

plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.grid()

plt.figure()

plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.grid()
plt.show()
```



Pablo Minango

- pablodavid218@gmail.com