

An Introduction to Deep Learning With Python

[8.2] DeepDream in Keras

Prof. Yuzo Iano

pgs: 281 - 286

Loading the pretrained Inception V3 model

```
In [4]: from keras.applications.inception_v3 import InceptionV3
        from keras import backend as K
```

```
In [6]: K.set_learning_phase(0)
        model = InceptionV3(weights = 'imagenet', include_top = False)
```

WARNING:tensorflow:From C:\Users\pablo\AppData\Roaming\Python\Python36\site-packages\tensorflow\python\framework\op_def_library.py:263: colocate_with (from tensorflow.python.framework.ops) is deprecated and will be removed in a future version.

Instructions for updating:

Colocations handled automatically by placer.

Downloading data from https://github.com/fchollet/deep-learning-models/releases/download/v0.5/inception_v3_weights_tf_dim_ordering_tf_kernels_notop.h5 (https://github.com/fchollet/deep-learning-models/releases/download/v0.5/inception_v3_weights_tf_dim_ordering_tf_kernels_notop.h5)

87916544/87910968 [=====] - 14s 0us/step

```
In [7]: layer_contributions = {
        'mixed2': 0.2,
        'mixed3': 3.,
        'mixed4': 2.,
        'mixed5': 1.5
        }
```

Defining the loss to be maximized

```
In [8]: layer_dict = dict([(layer.name, layer) for layer in model.layers])

        loss = K.variable(0.)
        for layer_name in layer_contributions:
            coeff = layer_contributions[layer_name]
            activation = layer_dict[layer_name].output

            scaling = K.prod(K.cast(K.shape(activation), 'float32'))
            loss += coeff * K.sum(K.square(activation[:, 2:-2, 2:-2, :])) / scaling
```

WARNING:tensorflow:Variable += will be deprecated. Use variable.assign_add if you want assignment to the variable value or 'x = x + y' if you want a new python Tensor object.

Gradient-ascent process

```
In [9]: dream = model.input
        grads = K.gradients(loss, dream)[0]
        grads /= K.maximum(K.mean(K.abs(grads)), 1e-7)
        outputs = [loss, grads]
        fetch_loss_and_grads = K.function([dream], outputs)
```

```
In [10]: def eval_loss_and_grads(x):
        outs = fetch_loss_and_grads([x])
        loss_value = outs[0]
        grad_values = outs[1]
        return loss_value, grad_values
```

```
In [11]: def gradient_ascent(x, iterations, step, max_loss=None):
        for i in range(iterations):
            loss_value, grad_values = eval_loss_and_grads(x)
            if max_loss is not None and loss_value > max_loss:
                break
            print('...Loss value at ', i, ':', loss_value)
            x += step * grad_values
        return x
```

Auxiliary functions

```
In [13]: import scipy
         from keras.preprocessing import image
```

```
In [14]: def resize_img(img, size):
         img = np.copy(img)
         factors = (1,
                    float(size[0]) / img.shape[1],
                    float(size[1]) / img.shape[2],
                    1)
         return scipy.ndimage.zoom(img, factors, order=1)
```

```
In [15]: def save_img(img, fname):
         pil_image = deprocess_image(np.copy(img))
         scipy.misc.imsave(fname, pil_image)
```

```
In [16]: def preprocess_image(image_path):
         img = image.load_img(image_path)
         img = image.img_to_array(img)
         img = np.expand_dims(img, axis=0)
         img = inception_v3.preprocess_input(img)
         return img
```

```
In [17]: def deprocess_image(x):
         if K.image_data_format() == 'channels_first':
             x = x.reshape((3, x.shape[2], x.shape[3]))
             x = x.transpose((1, 2, 0))
         else:
             x = x.reshape((x.shape[1], x.shape[2], 3))

         x /= 2.
         x += 0.5
         x *= 255.
         x = np.clip(x, 0, 255).astype('uint8')
         return x
```

Running gradient ascent over different successive scales

```
In [18]: import numpy as np
```

```
step = 0.01
num_octave = 3
octave_scale = 1.4
iterations = 20
max_loss = 10
base_image_path = 'grama.jpg'
img = preprocess_image(base_image_path)
original_shape = img.shape[1:3]
successive_shapes = [original_shape]

for i in range(1, num_octave):
    shape = tuple([int(dim / (octave_scale ** i))
                   for dim in original_shape])
    successive_shapes.append(shape)

successive_shapes = successive_shapes[::-1]

original_img = np.copy(img)
shrunk_original_img = resize_img(img, successive_shapes[0])

for shape in successive_shapes:
    print('Processing image shape', shape)
    img = resize_img(img, shape)
    img = gradient_ascent(img,
                          iterations=iterations,
                          step=step,
                          max_loss=max_loss)
    upscaled_shrunk_original_img = resize_img(shrunk_original_img, shape)
    same_size_original = resize_img(original_img, shape)
    lost_detail = same_size_original - upscaled_shrunk_original_img

    img += lost_detail
    shrunk_original_img = resize_img(original_img, shape)
    save_img(img, fname='dream_at_scales_' + str(shape) + '.png')

save_img(img, fname='final_dream.png')
```

```
Processing image shape (326, 326)
...Loss value at 0 : 1.2852011
...Loss value at 1 : 1.8864646
...Loss value at 2 : 2.5252922
...Loss value at 3 : 3.154289
...Loss value at 4 : 3.8254318
...Loss value at 5 : 4.517752
...Loss value at 6 : 5.1965256
...Loss value at 7 : 5.8220005
...Loss value at 8 : 6.468895
...Loss value at 9 : 7.0911365
...Loss value at 10 : 7.7676725
...Loss value at 11 : 8.401022
...Loss value at 12 : 9.07316
...Loss value at 13 : 9.727113
```

C:\Users\pablo\Python\envs\DAVID\lib\site-packages\ipykernel_launcher.py:3: DeprecationWarning: `imsave` is deprecated!

`imsave` is deprecated in SciPy 1.0.0, and will be removed in 1.2.0.

Use ``imageio.imwrite`` instead.

This is separate from the ipykernel package so we can avoid doing imports until

C:\Users\pablo\Python\envs\DAVID\lib\site-packages\scipy\ndimage\interpolation.py:605: UserWarning: From scipy 0.13.0, the output shape of zoom() is calculated with round() instead of int() - for these inputs the size of the returned array has changed.

"the returned array has changed.", UserWarning)

```
Processing image shape (457, 457)
...Loss value at 0 : 2.6484818
...Loss value at 1 : 3.9946883
...Loss value at 2 : 5.255889
...Loss value at 3 : 6.444748
...Loss value at 4 : 7.5588536
...Loss value at 5 : 8.785267
Processing image shape (640, 640)
...Loss value at 0 : 3.5263147
...Loss value at 1 : 6.0270605
...Loss value at 2 : 9.786003
```

```
In [21]: from IPython.display import Image
x = Image(filename='grama.jpg', height=400, width=400)
y = Image(filename='final_dream.png', height=400, width=400)
display(x, y)
```



```

In [22]: step = 0.01
num_octave = 3
octave_scale = 1.4
iterations = 20
max_loss = 10
base_image_path = 'ex2.jpg'
img = preprocess_image(base_image_path)
original_shape = img.shape[1:3]
successive_shapes = [original_shape]

for i in range(1, num_octave):
    shape = tuple([int(dim / (octave_scale ** i))
                   for dim in original_shape])
    successive_shapes.append(shape)

successive_shapes = successive_shapes[::-1]

original_img = np.copy(img)
shrunk_original_img = resize_img(img, successive_shapes[0])

for shape in successive_shapes:
    print('Processing image shape', shape)
    img = resize_img(img, shape)
    img = gradient_ascent(img,
                          iterations=iterations,
                          step=step,
                          max_loss=max_loss)
    upscaled_shrunk_original_img = resize_img(shrunk_original_img, shape)
    same_size_original = resize_img(original_img, shape)
    lost_detail = same_size_original - upscaled_shrunk_original_img

    img += lost_detail
    shrunk_original_img = resize_img(original_img, shape)
    save_img(img, fname='dream_at_scales_' + str(shape) + '.png')

save_img(img, fname='final_dream.png')

```

C:\Users\pablo\Python\envs\DAVID\lib\site-packages\scipy\ndimage\interpolation.py:605: UserWarning: From scipy 0.13.0, the output shape of zoom() is calculated with round() instead of int() - for these inputs the size of the returned array has changed.

"the returned array has changed.", UserWarning)

Processing image shape (234, 351)

```

...Loss value at 0 : 1.9095844
...Loss value at 1 : 2.4089239
...Loss value at 2 : 3.105619
...Loss value at 3 : 3.7721012
...Loss value at 4 : 4.5242815
...Loss value at 5 : 5.2202954
...Loss value at 6 : 5.8730936
...Loss value at 7 : 6.581296
...Loss value at 8 : 7.1658206
...Loss value at 9 : 7.7332554
...Loss value at 10 : 8.296839
...Loss value at 11 : 8.777473
...Loss value at 12 : 9.223952
...Loss value at 13 : 9.648778

```

C:\Users\pablo\Python\envs\DAVID\lib\site-packages\ipykernel_launcher.py:3: DeprecationWarning: `imsave` is deprecated!

`imsave` is deprecated in SciPy 1.0.0, and will be removed in 1.2.0.

Use ``imageio.imwrite`` instead.

This is separate from the ipykernel package so we can avoid doing imports until

Processing image shape (327, 491)

```

...Loss value at 0 : 3.0115323
...Loss value at 1 : 4.412323
...Loss value at 2 : 5.4962444
...Loss value at 3 : 6.409892
...Loss value at 4 : 7.2163553
...Loss value at 5 : 7.989553
...Loss value at 6 : 8.698456
...Loss value at 7 : 9.352154
...Loss value at 8 : 9.979662

```


Processing image shape (459, 688)

```

...Loss value at 0 : 3.0647936
...Loss value at 1 : 4.3328824
...Loss value at 2 : 5.38324
...Loss value at 3 : 6.3031907
...Loss value at 4 : 7.191517
...Loss value at 5 : 7.9933796
...Loss value at 6 : 8.735939
...Loss value at 7 : 9.444316

```

```
In [23]: from IPython.display import Image
x = Image(filename='ex2.jpg', height=400, width=400)
y = Image(filename='final_dream.png', height=400, width=400)
display(x, y)
```



Pablo Minango

- pablodavid218@gmail.com (<mailto:pablodavid218@gmail.com>)