

# An Introduction to Deep Learning With Python

## [8.4] Generating images with variational autoencoders

Prof. Yuzo Iano

pgs: 300 - 303

### VAE encoder network

```
In [1]: import keras
        from keras import layers
        from keras import backend as K
        from keras.models import Model
        import numpy as np
```

Using TensorFlow backend.

```
In [2]: img_shape = (28, 28, 1)
        batch_size = 16
        latent_dim = 2 # Dimensionality of the latent space: a plane

        input_img = keras.Input(shape=img_shape)

        x = layers.Conv2D(32, 3,
                          padding='same', activation='relu')(input_img)
        x = layers.Conv2D(64, 3,
                          padding='same', activation='relu',
                          strides=(2, 2))(x)
        x = layers.Conv2D(64, 3,
                          padding='same', activation='relu')(x)
        x = layers.Conv2D(64, 3,
                          padding='same', activation='relu')(x)
        shape_before_flattening = K.int_shape(x)

        x = layers.Flatten()(x)
        x = layers.Dense(32, activation='relu')(x)

        z_mean = layers.Dense(latent_dim)(x)
        z_log_var = layers.Dense(latent_dim)(x)
```

WARNING:tensorflow:From C:\Users\pablo\AppData\Roaming\Python\Python36\site-packages\tensorflow\python\framework\ops\_def\_library.py:263: colocate\_with (from tensorflow.python.framework.ops) is deprecated and will be removed in a future version.

Instructions for updating:  
Colocations handled automatically by placer.

### Latent-space-sampling function

```
In [3]: def sampling(args):
        z_mean, z_log_var = args
        epsilon = K.random_normal(shape=(K.shape(z_mean)[0], latent_dim),
                                   mean=0., stddev=1.)
        return z_mean + K.exp(z_log_var) * epsilon
```

```
In [4]: z = layers.Lambda(sampling)([z_mean, z_log_var])
```

### VAE decoder network, mapping latent space points to images

```
In [5]: decoder_input = layers.Input(K.int_shape(z)[1:])

x = layers.Dense(np.prod(shape_before_flattening[1:]),
                  activation='relu')(decoder_input)
x = layers.Reshape(shape_before_flattening[1:])(x)

x = layers.Conv2DTranspose(32, 3,
                           padding='same', activation='relu',
                           strides=(2, 2))(x)
x = layers.Conv2D(1, 3,
                  padding='same', activation='sigmoid')(x)

decoder = Model(decoder_input, x)

z_decoded = decoder(z)
```

#### Custom layer used to compute the VAE loss

```
In [6]: class CustomVariationalLayer(keras.layers.Layer):

    def vae_loss(self, x, z_decoded):
        x = K.flatten(x)
        z_decoded = K.flatten(z_decoded)
        xent_loss = keras.metrics.binary_crossentropy(x, z_decoded)
        kl_loss = -5e-4 * K.mean(
            1 + z_log_var - K.square(z_mean) - K.exp(z_log_var), axis=-1)
        return K.mean(xent_loss + kl_loss)

    def call(self, inputs):
        x = inputs[0]
        z_decoded = inputs[1]
        loss = self.vae_loss(x, z_decoded)
        self.add_loss(loss, inputs=inputs)
        return x
```

```
In [7]: y = CustomVariationalLayer()([input_img, z_decoded])
```

#### Training the VAE

```
In [8]: from keras.datasets import mnist

vae = Model(input_img, y)
vae.compile(optimizer='rmsprop', loss=None)
vae.summary()
```

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	(None, 28, 28, 1)	0	
conv2d_1 (Conv2D)	(None, 28, 28, 32)	320	input_1[0][0]
conv2d_2 (Conv2D)	(None, 14, 14, 64)	18496	conv2d_1[0][0]
conv2d_3 (Conv2D)	(None, 14, 14, 64)	36928	conv2d_2[0][0]
conv2d_4 (Conv2D)	(None, 14, 14, 64)	36928	conv2d_3[0][0]
flatten_1 (Flatten)	(None, 12544)	0	conv2d_4[0][0]
dense_1 (Dense)	(None, 32)	401440	flatten_1[0][0]
dense_2 (Dense)	(None, 2)	66	dense_1[0][0]
dense_3 (Dense)	(None, 2)	66	dense_1[0][0]
lambda_1 (Lambda)	(None, 2)	0	dense_2[0][0] dense_3[0][0]
model_1 (Model)	(None, 28, 28, 1)	56385	lambda_1[0][0]
custom_variational_layer_1 (Custom Layer)	(None, 28, 28, 1)	0	input_1[0][0] model_1[1][0]
Total params: 550,629			
Trainable params: 550,629			
Non-trainable params: 0			

```
In [9]: (x_train, _), (x_test, y_test) = mnist.load_data()
```

```
x_train = x_train.astype('float32') / 255.  
x_train = x_train.reshape(x_train.shape + (1,))  
x_test = x_test.astype('float32') / 255.  
x_test = x_test.reshape(x_test.shape + (1,))
```

```
vae.fit(x=x_train, y=None,  
       shuffle=True,  
       epochs=10,  
       batch_size=batch_size,  
       validation_data=(x_test, None))
```

WARNING:tensorflow:From C:\Users\pablo\AppData\Roaming\Python\Python36\site-packages\tensorflow\python\ops\math\_ops.py:3066: to\_int32 (from tensorflow.python.ops.math\_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.cast instead.

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/10  
60000/60000 [=====] - 314s 5ms/step - loss: 2.2031 - val_loss: 0.1994  
Epoch 2/10  
60000/60000 [=====] - 455s 8ms/step - loss: 0.1955 - val_loss: 0.1920  
Epoch 3/10  
60000/60000 [=====] - 447s 7ms/step - loss: 0.1911 - val_loss: 0.1893  
Epoch 4/10  
60000/60000 [=====] - 428s 7ms/step - loss: 0.1887 - val_loss: 0.1875  
Epoch 5/10  
60000/60000 [=====] - 426s 7ms/step - loss: 0.1870 - val_loss: 0.1864  
Epoch 6/10  
60000/60000 [=====] - 428s 7ms/step - loss: 0.1857 - val_loss: 0.1856  
Epoch 7/10  
60000/60000 [=====] - 430s 7ms/step - loss: 0.1847 - val_loss: 0.1847  
Epoch 8/10  
60000/60000 [=====] - 402s 7ms/step - loss: 0.1839 - val_loss: 0.1834  
Epoch 9/10  
60000/60000 [=====] - 402s 7ms/step - loss: 0.1832 - val_loss: 0.1823  
Epoch 10/10  
60000/60000 [=====] - 401s 7ms/step - loss: 0.1828 - val_loss: 0.1833
```

```
Out[9]: <keras.callbacks.History at 0x1f921503c18>
```

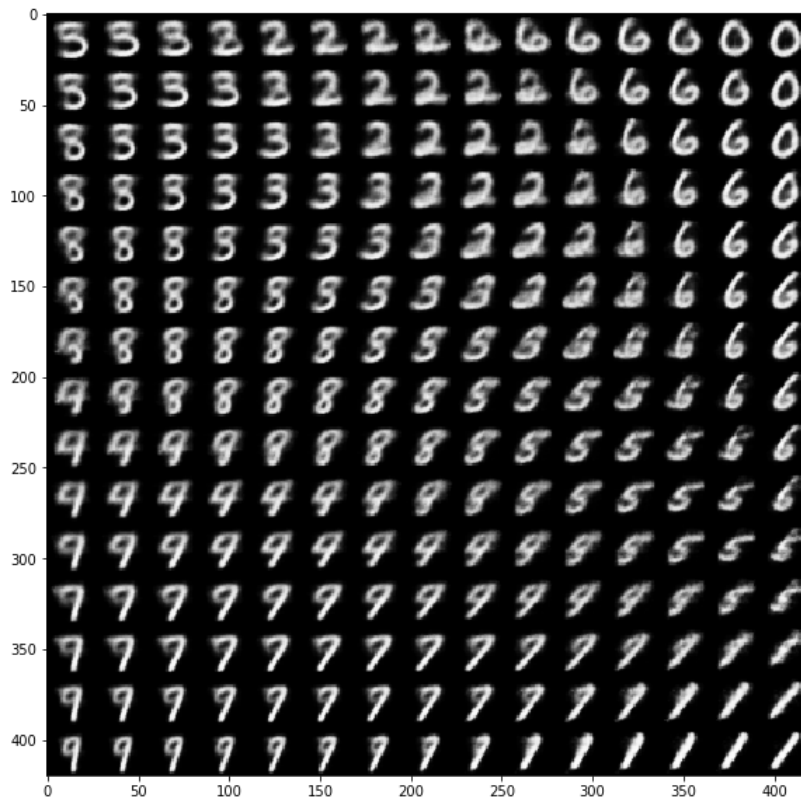
**Sampling a grid of points from the 2D latent space and decoding them to images**

```
In [11]: import matplotlib.pyplot as plt
from scipy.stats import norm

n = 15
digit_size = 28
figure = np.zeros((digit_size * n, digit_size * n))
grid_x = norm.ppf(np.linspace(0.05, 0.95, n))
grid_y = norm.ppf(np.linspace(0.05, 0.95, n))

for i, yi in enumerate(grid_x):
    for j, xi in enumerate(grid_y):
        z_sample = np.array([[xi, yi]])
        z_sample = np.tile(z_sample, batch_size).reshape(batch_size, 2)
        x_decoded = decoder.predict(z_sample, batch_size=batch_size)
        digit = x_decoded[0].reshape(digit_size, digit_size)
        figure[i * digit_size: (i + 1) * digit_size,
              j * digit_size: (j + 1) * digit_size] = digit

plt.figure(figsize=(10, 10))
plt.imshow(figure, cmap='Greys_r')
plt.show()
```



**Pablo Minango**

- pablodavid218@gmail.com